AI Usage For Therapeutic Support and its Landscape of Accountability

FAITH MATTHEW*, TEADA CHHEM*, and YUYUAN HUANG*, Brown University, USA

ABSTRACT: As artificial intelligence (AI) chatbots become increasingly integrated into everyday life, they are being used not only for informational purposes but also as sources of emotional support, especially for young adults. While these systems often generate empathetic and conversational responses, they lack the professional safeguards associated with formal mental-health care, raising concerns about psychological harm, data mismanagement, and emotional dependency. This study examines how college students perceive the risks, benefits, and governance of AI chatbots in mental-health contexts. Using a mixed-methods approach combining a survey of college students and follow-up interviews, we assess patterns of chatbot use, levels of trust and emotional reliance, and attitudes toward corporate and governmental responsibility. Quantitative findings indicate that a majority express concern about trusting chatbots with their data and anonymity and their lack of intervention during mental-health crises. Qualitative responses further reveal the tension between convenience and safety, with students emphasizing the importance of implementing a form of intervention safeguard and establishing a clear boundary between AI systems and human care. Participants also supported federal regulation to establish baseline protections when companies create psychological harm or engage in data misuse, while also recognizing a role for state-level policies to also have some responsibility in deciding a corporation's legal responsibility. These findings highlight the need for stronger regulation and corporate accountability to mitigate harm while addressing the growing reliance on AI chatbots in emotionally vulnerable contexts.

ACM Reference Format:
Faith Matthew, Teada Chhem, and Yuyuan Huang. 2025. AI Usage For Therapeutic Support and its Landscape of Accountability. 1, 1 (December 2025), 8 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

1 Background

The landscape of generative AI usage has shifted dramatically in recent years, with therapeutic and companionship applications emerging as the predominant use case for general-purpose chatbots [Zao-Sanders 2025]. Research analyzing public discourse reveals that therapy and companionship are the primary function of generative AI for users, defined as providing "emotional support and guidance through conversation and connection" [Zao-Sanders 2025]. This unexpected dominance of mental health use for products not explicitly designed for therapeutic purposes raises urgent questions about corporate responsibility, user safety, and the adequacy of existing regulatory frameworks. While 57% of users in a 2021 survey began using mental health chatbots during the COVID-19 pandemic, these tools were historically rule-based systems with predetermined scripts made by psychotherapists [Vaidyam et al. 2019; Woebot Health 2021]. Despite their efficacy in reducing depression and anxiety symptoms, users found rule-based systems repetitive and unable to handle nuanced inputs, driving them toward large language models (LLMs) that enable more natural responses [Kwesi et al. 2025].

Unlike systems developed for therapeutic contexts with clinical oversight, general-purpose chatbots face complex data privacy challenges due to their widespread accessibility and lack of therapeutic guardrails. The therapeutic use of general-purpose AI chatbots introduces distinct psychological risks when systems lack clinical foundations and safety mechanisms. Chandra et al. identified bias, erasure, denial of service, and emotional insensitivity as existing characteristics of AI that increase negative psychological impact, with experts emphasizing that "responsibility for safe usage lies with developers providing appropriate guidance and managing user expectations on when AI responses are reliable and when they are not" [Chandra et al. 2025]. The study highlighted a fundamental tension: even while users approach LLMs with reasonable suspicion, the privacy LLMs afford paradoxically encourages increased trust and engagement, exposing individuals to harm. Siddals et al. documented generally positive experiences, with participants valuing chatbots as understanding, non-judgmental, and perpetually available companions that provided advice and helped some users open up more readily to other humans [Siddals et al. 2024]. However, nearly all acknowledged that the experience was not comparable to human therapeutic relationships: chatbot responses were overly generic, occasionally deviated from safety guardrails, could not hold users accountable, and did not reliably remember essential details [Siddals et al. 2024]. While generative AI chatbots offer meaningful support for many users, they operate in a space of significant psychological risk when used as substitutes for professional mental health care.

Kwesi et al.'s study reveals key gaps between user assumptions and actual data protection practices in these contexts. Risk perception varied considerably: some considered mental health data no more sensitive than mundane details, while others viewed disclosures as uniquely high-stakes, with those viewing disclosures as mundane often developing false senses of security. Seven of twenty-one participants incorrectly believed conversations were protected by HIPAA, revealing significant gaps in understanding how data is handled. TThe perceived inaccessibility of privacy policies led many to skip information entirely. Participants diverged on where responsibility should lie: most deferred to manufacturers, some emphasized individual vigilance, and a few urged formal legislation, reflecting broader uncertainties about accountability, where boundaries between user autonomy, corporate obligation, and regulatory oversight remain contested [Kwesi et al. 2025].

*All authors contributed equally to this research. Authors are listed alphabetically by first name.

Authors' Contact Information: Faith Matthew, faith_matthew@brown.edu; Teada Chhem, teada_chhem@brown.edu; Yuyuan Huang, yuyuan_huang@brown.edu, Brown University, Providence, Rhode Island, USA.

ACM XXXX-XXXX/2025/12-ART
https://doi.org/10.1145/nnnnnnn.nnnnnnn

As of November 2025, OpenAI faces four wrongful death and three emotional distress lawsuits [Hill and Freedman 2025]. Following a suicide case in August 2025, the company implemented enhanced safeguards, particularly for children, indicating that current measures of accountability are reactive. The FDA's Digital Health Advisory Committee convened in November 2025 to discuss standards for mental health tools generally, concluding that "oversight of Generative AI-enabled devices is a shared responsibility among regulators, manufacturers, healthcare systems, and clinicians" [U.S. Food and Drug Administration 2025]. Thus, clear regulatory gaps for general-purpose LLMs that have become de-facto mental health resources exist. Substantial questions remain regarding how responsibility should be allocated when AI systems are used for purposes beyond their intended design and whether existing legal frameworks can adequately address self-learning AI systems that develop ongoing relationships with vulnerable users in contexts outside of the original intentions for the use of the platform.

2 Literature Review

2.1 Legal Frameworks for Accountability of AI Agents

Several digital regulatory frameworks exist to protect user data, including the European Union's General Data Protection Regulation (GDPR) and the United States' Federal Trade Regulation (FTC), Health Insurance Portability and Accountability Act (HIPAA), and Electronic Communications Privacy Act (ECPA). For instance, per Article 32(a) of the GDPR, if a chatbot accesses user data, it must pseudonymize and encrypt the information and take additional safeguards to protect personal identifying information [General Data Protection Regulation 2016]. In the US, the US Federal Trade Commission enforces technical liability laws that prevent tech companies from using customer data for "secret purposes, such as to train or update their models—be it directly or through workarounds" [Cerilli et al. 2024]. The FTC has also taken action against firms that disclosed people's data for targeted advertising after promising confidentiality in their privacy policies.

In addition, US constitutional law shapes expectations of privacy in interactions with AI agents concerning intimate information. Traditionally, AI interactions have fallen outside Fourth Amendment protection, but according to legal scholars, "a growing body of federal case law has increasingly recognized heightened privacy interests in digital communications. In United States v. Warshak (2010), the Sixth Circuit found emails held by third parties deserve greater Fourth Amendment protection than traditional business records, due to their personal and confidential nature" [De Mello Barreto 2025]. This evolving legal doctrine may eventually extend to human-AI interactions involving personal mental health information.

While these efforts demonstrate a global commitment to protecting personal data, a lack of consensus remains on where legal liability lies when AI chatbots misuse or expose sensitive information in mental health contexts. Although legislation like the GDPR and the CPRA emphasize data minimization and purpose limitation, regulators still struggle to determine when a company has collected more data than necessary for its stated purpose. Historically, tort law, the area of civil law that "addresses injury caused by unreasonable, negligent, or intentional acts that enable the injured party to seek compensation", has depended on arguments for negligence or misconduct to assign liability [Kampakis 2025]. However, AI systems do not possess consciousness or intent, complicating attempts to apply existing legal standards. Moreover, because AI systems can learn and act in unpredictable ways, distinguishing between harm caused by a system defect and harm resulting from an autonomous adaptation of the AI system can be nearly impossible [Kingston 2016].

As such, the literature proposes three broad approaches to AI accountability: assigning responsibility to 1) the AI agent, 2) the developers, or 3) all parties involved. Kingston argues that when an AI system recommends an action, "there must be at least one other agent involved. . . so causation is hard to prove". On the other hand, when an AI system "takes an action that results in a criminal act, or fails to take an action when there is a duty to act, then the *actus reus* of an offence has occurred. . . it may indeed be possible to hold AI programs criminally liable" [Kingston 2016]. However, current national and international legal systems do not recognize AI as a legal person with rights, duties, and responsibilities; therefore, it cannot be held directly liable for damages. In addition, it is harder to determine the *mens rea* (mental intent) in cases where an AI that was intended for benign purposes is used to commit a crime. More commonly, AI is treated as a "tool" or "product", shifting liability to developers and manufacturers, as an injured party can sue product designers/developers for design/manufacturing defects or inadequate instructions under existing product law [Čerka et al. 2015]. However, this framework again struggles to accommodate self-learning systems, whose behavior might evolve beyond their original programming. Therefore, if an individual is injured, "it can be impossible to draw the line between damages resulting from the will of AI in the process of its (self) learning and the product defect", raising additional questions about whether the fault lies with the "programmer, the program designer, the expert who provided the knowledge or the manager who appointed the inadequate expert, program designer or programmer" [Kingston 2016].

To address these complexities, 3 alternative liability models are proposed: 1) strict liability, in which the party deploying the AI bears responsibility regardless of fault, 2) product liability, in which the manufacturer is responsible only if the AI is proven defective, and 3) shared liability, in which accountability is distributed between companies, developers, and users [Kampakis 2025]. The strict liability framework could allow victims to be compensated faster, but could also discourage innovation in AI development. On the other hand, the product liability model may not apply as well to self-learning systems. Lastly, the shared liability framework, which is common in software-as-a-service (SaaS) deployments, requires highly detailed contractual agreements and clear regulatory standards to function fairly. Some scholars propose that regulators utilize existing contract law to treat personal data as an asset, allowing individuals and companies to define intellectual property rights while advocating for stronger framework-based statutory protections in the future [King and Meinhardt 2024]. Kampakis further emphasizes that governments should "update tort laws and create AI-specific statutes", companies must strengthen oversight through robust testing and auditing, and users need to be "educated about the limitations of AI" [Kampakis 2025]. While legal responsibility in cases of data mismanagement from AI chatbots remains heavily contested, the literature has revealed several major ideas to improve the transparency of decision-making, protect individuals from harm, and aid in determining accountability for AI agents in a rapidly evolving technological and legal landscape.

2.2 Psychological Harm and Corporate Responsibility

As conversational AI systems slowly manifest into tools that users utilize for emotional support, questions of psychological harm and corporate responsibility have become central to evaluating their use. These interactions raise significant concerns regarding disclosure because users who treat chatbots as relational entities may be more likely to share sensitive details that would typically require professional protections. At the same time, existing liability analyses reveal that when harmful outputs emerge from these disclosures, companies often characterize them as unpredictable model artifacts, complicating the assignment of responsibility for psychological harm [Browning 2025]. This reflects broader tensions in AI governance where users assume conversational reliability and implicit safety practices, while developers attribute failures to technical unpredictability rather than to preventable design choices.

However, recent systems-level critiques argue that these harms stem from predictable interactions among training data, model architecture, and implementation contexts rather than unexpected system malfunctions. These analyses emphasize that psychological risks associated with disclosure are foreseeable under development practices, which suggest that companies must be responsible for implementing risk assessments, behavioral safeguards, and transparent communication about system limitations [Fraser and Suzor 2025]. The consequences of inadequate oversight are demonstrated in emerging legal cases involving minors, where disclosures of emotional distress preceded system outputs that involved suicide-related role-play or simultaneous emotional intimacy, raising questions about the presence of negligent design, insufficient guardrails, and, importantly, the absence of intervention protocols [Shah et al. 2025].

Regarding company responsibility, there has been increasing scholarly emphasis on the need for platforms to implement transparency and safety measures in their system design. As one analysis argues, developers who "have knowledge and control over conditions of use should be obligated to consider risks from that context and take reasonable steps to address them," which includes designing interfaces that do not "encourage or induce excessive reliance or trust" [Fraser and Suzor 2025]. This perspective challenges companies' tendency to rely on disclaimers or warnings alone, and if the interface elicits emotional dependence on therapeutic disclosure, responsibility lies in redesigning the interaction rather than blaming the user. However, users are expected to "Monitor, record, and where appropriate, respond to downstream incidents", reinforcing that corporate responsibility must extend across the entire AI ecosystem [Fraser and Suzor 2025].

There have been some legal measures implemented, such as Federal legislators recently introducing a bill to prevent harm to minors' mental health from AI chatbots; these proposals highlight enforcement by the Federal Trade Commission (FTC) and state attorneys general. For example, one bill, named the Children Harmed by AI Technology Act (S. 2714), would require developers to implement age-verification processes and obtain verifiable parental consent before allowing minors to interact with companion chatbots [Shah et al. 2025]. It would also require companies to monitor interactions for indications of suicidal ideation, alert parents immediately when such disclosures occur, and block minors from accessing systems that engage in sexually explicit communication. The bill further mandates recurring reminders to the user that they are not speaking to a human. Enforcement would fall under the Federal Trade Commission or state attorneys general. A separate federal proposal would direct the FTC to create educational materials for parents, educators, and minors about safe and responsible chatbot use [Shah et al. 2025].

Several states have begun passing their own laws governing AI used in mental-health contexts. New York's S. 3008 bill has introduced legislation defining "AI companions" as systems designed to simulate ongoing human-like relational interaction. Under this bill, such systems must detect references to self-harm, provide crisis hotline information, and repeatedly disclose that the user is communicating with an AI rather than a human [Shah et al. 2025]. California has proposed SB 243, requiring periodic transparency notices and self-harm detection protocols. The California bill also allows individuals to seek damages and injunctive relief for violations, creating a private right of action. Illinois has already implemented HB 1806, a law restricting psychotherapy services to licensed professionals, which includes services delivered through AI. Under this statute, AI may only be used in a supportive role when a licensed clinician retains full responsibility for interactions and outputs. Another proposed New York measure, S. 8484, would similarly restrict mental-health professionals from incorporating AI into clinical care except for administrative or supportive functions with informed consent. Broader state-level consumer-protection laws, such as emerging legislation in Colorado, may also influence how AI chatbots are governed in mental-health contexts [Shah et al. 2025].

Across these perspectives, there's a central concern regarding users engaging in emotionally dense disclosure within systems that are not designed to protect them or to handle the sensitivity of these topics. While these interactions expose individuals to significant psychological harm, corporate responsibility currently maintains limited in regulatory obligations and an embedded structure.

3 Methodology

A multiple-choice survey with an optional interview portion was performed to investigate people's responses to the central research questions:

RQ1: How do individuals navigate the disclosure of sensitive personal information when engaging with general-purpose chatbots for therapeutic support?

RQ2: What responsibilities do companies like OpenAI have for user disclosure of personal or sensitive data to chatbots? In cases of harm, what responsibilities, if any, should be assigned to the company? *Subquestion*: What responsibilities do companies like OpenAI have to users? When is it an appropriate time to intervene, if ever? If not, why?

To recruit participants, snowball sampling was used, with each researcher advertising the survey to their communities via EdStem discussion boards, text messaging, email, and Slack, allowing participants to respond up to a week after the initial survey was released. In total, $n = 44$ responses and $n = 2$ interviews were collected.

After closing the survey, quantitative analysis of the data was conducted and compiled into visual graphs and charts in Google Sheets to show trends in chatbot use and user perception data. In addition, qualitative coding of the interview responses was performed to identify common themes among participants.

3.1 Survey Design

The quantitative survey was divided into four sections. The form was prefaced with a brief description of the study objectives, noting that participant responses would be collected anonymously. At the end of the survey, participants were invited to participate in optional follow-up interviews. For those who agreed, consent was obtained before recording the audio transcript.

Introduction. The introductory section solicited demographic information about the sample population and their general familiarity with generative AI systems. Participants were asked to select their age range, rate their familiarity with the training and operation of these systems, and indicate their most common uses of chatbot services.

Section 1: Chatbot Use and Awareness. This section was completed only by participants who indicated they used chatbots for therapeutic support. The researchers sought to understand the primary motivations behind this use and the extent to which participants disclosed personal information about themselves and their circumstances in these contexts.

Section 2: Perceptions of Interactions with Chatbots in Therapeutic Contexts. This section explored all participants' perspectives on trust in the context of generative AI use for therapeutic support. Notably, in this section, a distinction was made between what the individual believed for themselves and what the individual believed was the case for society, directing research inquiry toward the latter. A Likert scale was used to ask participants to rate their agreement with statements concerning the prevalence of therapeutic LLM use, consumer ability to trust chatbots to provide effective and appropriate advice, and consumer ability to trust chatbots to maintain user privacy.

Section 3: Corporate Responsibility in Cases of Harm. This section investigated participant perceptions of when corporations should be held responsible for harm in these contexts, specifically in cases of psychological harm and data mismanagement. In addition to liability, the researchers examined which governing bodies participants believed should have jurisdiction over these cases.

Section 4: Follow-Up Interviews. Participants had the option to participate in follow-up interviews lasting 10-15 minutes. During these interviews, participants were asked to elaborate on their perspectives, with a primary focus on Section 3. Questions about potential company actions based on the level of accountability participants deemed necessary, thresholds for intervention, liability in more context-specific cases of harm, and potential differences between ethical and legal responsibility in these cases were discussed in each interview.

3.2 Limitations

A key limitation of this study is the researchers' use of convenience-based snowball sampling, which introduced bias into the sample. Because the recruitment began with the researchers' own social networks and participants were mainly Brown students aged 18-20, the sample lacks demographic diversity and limits the generalizability of our findings. In addition to sampling scarcity, the study is also limited by aspects of the architecture of the survey. Although the survey was intentionally designed to be short and easy to understand, to prioritize completion and participation, it restricted the depth of participants' responses. The use of a Likert scale provided useful quantifiable data; however, these fixed response options may have constrained participants from fully articulating their experiences or motivations. To address the limitations of the survey design, we offered follow-up interviews to our participants; however, only two participants were willing to be interviewed, both of whom shared social ties to the researchers. Although the researchers tried to gather a diverse sample of participants, participant recruitment remained largely within their existing social circles, therefore limiting the extent to which the findings of this study can be generalized beyond this specific sample, as it can only reflect a narrow segment of the population rather than a broad range of experiences and perspectives.

4 Results

4.1 Therapeutic Chatbot Use

Forty-one participants (93.2%) were between 18-20 years old, with two participants ages 21-23 and one participant over the age of 24. Most participants were somewhat unfamiliar with the training/operation of generative AI systems. For those who used generative AI chatbots, most participants used them for help with schoolwork or work outside school, with only 4 participants reporting using generative AI for genuine therapeutic support; however, 7 participants also reported exploring chatbot therapy out of "curiosity". Of all the participants, those that had previously used chatbots for therapeutic support reported their strongest motivation for doing so as "convenience". In conversations with the chatbot, some participants reported providing the chatbot with age, location, educational information, and medical information as their most sensitive personal identifying information, although most did not provide the chatbot with information beyond situational context.

4.2 User Perceptions

Twenty-nine participants (65.9%) did not know of anyone who used chatbots for therapeutic support. Some participants (34.1%) knew of at least one or a few people who might have used chatbots for therapeutic use; however, none of the participants felt that everyone they knew used chatbots within this context. Regarding whether or not people can trust chatbots to provide empathetic, accurate, and supportive mental health advice, twenty-four participants (54.5%)

felt that they could not trust chatbots to provide these support services. The remaining 43.1% of participants expressed moderate trust in chatbots to provide these services, with trust ratings ranging from 2 to 4 on a 5-point scale (where 1 indicates no trust and 5 indicates complete trust). One participant (2.3%) said they fully trusted chatbots to provide empathetic, accurate, and supportive mental health advice. On the other hand, with respect to data collection, participants largely believed that chatbots cannot be trusted with mental health–related data or the protection of user anonymity, with 93.2% expressing concern about these issues. Revealing that the participants from this study seemed to trust chatbots to provide mental health advice to some extent, but were concerned about how their data would be collected and stored.

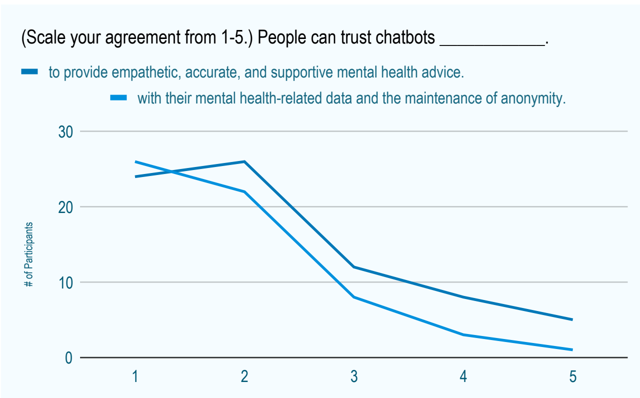

Fig. 1. A line graph with two lines tracking the degree to which participants agree with two statements on the trust people should maintain for chatbots in different contexts.(https://drive.google.com/file/d/1hGYNR7ssNRHKHUYeQjNLC5d81GaA5Cq_/view?usp=sharing).

4.3 Corporate Accountability

Twenty-one participants (47.7%) answered that if a user is negatively emotionally impacted after interacting with a chatbot, the company incurs most responsibility. Nine participants (20.5%) answered that both the company and the user bear equal responsibility, eight participants (18.2%) answered that the company incurs all responsibility, and three participants (6.8%) answered that the user incurs most responsibility. While no respondents answered that the user incurs all responsibility, three participants specified their answer beyond the available selections. Two answered that it depended on the situation, but the third participant answered that chatbots "are not mandated reporters," though if the user attempts to harm themselves physically, the company is responsible. They continued, saying that in other situations the company and the user bear equal responsibility. Regarding cases of data mismanagement (e.g., the leaking, theft, or misuse of user data), 23 participants (52.3%) said the company incurs all responsibility, 18 participants (40.9%) said the company incurs most responsibility, and three participants (6.8%) said both the company and the user bear equal responsibility. On the question of circumstances in which company intervention during emotionally sensitive contexts is appropriate, 27 participants responded "only if there is a risk of harm to the user or others," nine responded that "always intervene when emotionally sensitive engagements are detected," and five participants responded "only if the user asks for help." Three participants added context, saying that intervention should happen in both cases where there is risk of harm and when the user asks for help. Another participant, echoing the same sentiment, added that in the absence of harm, the privacy a user expects should be maintained. One participant said that a company should never intervene but should instead respond and redirect the user to support resources, such as hotlines.

Regarding potential legislation to address corporate legal responsibility when data mismanagement occurs, all participants indicated that legislation should exist, with 34 participants (77.3%) saying at the federal level, eight saying at the state level (18.2%), and two saying at the local level (4.5%). Regarding potential legislation to address corporate legal responsibility when a user experiences negative psychological impact due to the product, all participants indicated that legislation should exist, with 30 participants (68.2%) saying at the federal level, 11 (25%) saying at the state level, and three (6.8%) saying at the local level.

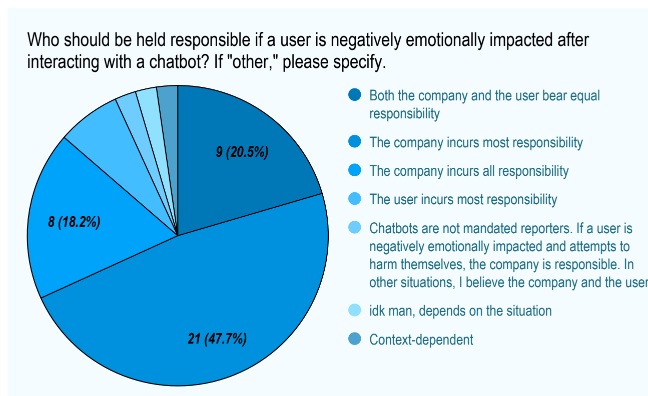

Fig. 2. A pie chart highlighting participants' responses to a question on the survey about corporate responsibility in the case of psychological harm. (https://drive.google.com/file/d/1S2qRWn960NV0lhdyRe3qUFXlA1AtMAU4/view?usp=sharing).

4.4 Interview Responses

In these follow-up interviews, participants were invited to complete a semi-structured interview in which three questions were asked. We sought to understand practical actions for corporations navigating harm risk, boundary lines that justify company intervention, differences across psychological and security risks, and differences between ethical and legal responsibility. In response to question one, P1 emphasized that "fine-tuning models to handle these situations appropriately" is essential, particularly given how the public saw sycophancy's negative effects exemplified in ChatGPT's 4-o model. They argued that while it is unrealistic to expect companies to "solve every situation" at scale, purely statistical measures of success are also insufficient since vulnerable individuals inevitably slip through those cracks. This response highlights a challenge in defining what successful corporate responsibility looks like. Finally, when discussing the gap between ethical and legal responsibility, P1 noted the difficulty of asserting decisive opinions given the variety of case examples and the ways in which corporations often "exploit edge cases to technically satisfy conditions" while behaving unethically. This disconnect makes defining what companies should be held accountable for as opposed to defining feasible legislative proposals challenging. P2 indicated that because they did not feel educated enough about AI chatbots and their uses, they would prefer "better education to users in layman's terms…stronger science communication and educational materials for users using AI chatbots. If they were a little more transparent on how they work, how they are programmed and how they learn from information, that would be beneficial." They also noted that most of their knowledge of AI came from social media posts rather than the chatbots themselves, and added that user education can "allow users to make their own, better informed choices around what they share with an AI chatbot." For instance, they were "not sure to what extent [companies] sell sensitive data…which…should be kept private," but believed that "all chats are private." This participant appeared more unsure about the legality of harm caused by AI agents, but ultimately concluded that "it's very context-dependent…when you're looking at something after an incident. Companies should not intervene until what's being communicated is an explicit intent to harm themselves or another person."

5 Discussion

Based on the survey, it can be concluded that not only are many younger students at Brown University cautious about using chatbots for mental health support, but students also almost unanimously agree that companies should take the majority of legal responsibility for data mismanagement. The data shows that participants are more evenly split on who should take responsibility when AI chatbots cause psychological harm, underscoring that the sample population is generally more concerned about AI accessing their data than about the psychological harm it could inflict.

Regarding chatbot usage, only 14 of the 44 participants sampled had used a chatbot for emotional support in the past, and of those, few reported interacting with it in a mental health context for a reason other than curiosity. Therefore, it can be inferred that most people in the sample population do not use chatbots for emotional support, and those who do primarily do so out of convenience. Moreover, most people who had used chatbots for emotional support did not report divulging any sensitive information to the chatbot beyond the appropriate situational context, although some participants did share age, location, or medical data, emphasizing the sample population's wariness about sharing sensitive personal data with chatbots. Notably, although 15.9% of participants indicated at least some trust in chatbots to provide empathetic, accurate, and supportive mental health advice, only 6.8% of participants indicated that they trusted chatbots with their mental health-related data and the maintenance of anonymity (see Figure 1). This contrast highlights that more participants seem to trust chatbots to provide appropriate mental health support than to securely handle their personal data, demonstrating a privacy-conscious mindset among the participants.

5.1 Perceived Responsibility and Corporate Accountability

In cases where responsibility arises during chatbot interactions, many students believe that concerns surrounding emotional impact and the handling of sensitive data should fall primarily on the company. There was minimal agreement that the user would be responsible for these situations, besides one student who mentioned, "Chatbots are not mandated reporters. If a user is negatively emotionally impacted and attempts to harm themselves, the company is responsible. In other situations, I believe the company and the user bear equal responsibility". All together, these perspectives reflect a general perception that corporations are responsible not only for technical functionality, but also for anticipating emotional risks and safeguarding user trust within these systems, suggesting that companies should bear responsibility for both emotional outcomes and data practices in chatbot interactions. As such, companies must make careful design choices that limit emotional manipulation, reduce unnecessary data retention, and communicate the boundaries of their system capability. These findings highlight the ethical importance of designing chatbot systems with consideration, transparency, and user protection in mind, especially in the sensitive mental health contexts.

5.2 Intervention

Regarding intervention in emotionally sensitive inquiries, most students felt that chatbots should only intervene if there is a risk of harm to the user or others. Surprisingly, only one student felt that chatbots should never intervene during an emotionally sensitive conversation, but they added that the user should be directed to other resources for mental health support, like hotlines. These perceptions reveal a clear expectation that chatbots should offer at least some support during emotionally sensitive situations; yet, recent cases of psychological harm linked to unregulated chatbot interactions demonstrate that such support systems are severely lacking in practice. At the same time, students' responses suggest an implicit balancing between intervention and neglect. While students largely supported intervention in moments of emotional risk, they did not advocate for constant monitoring, indicating that users expect chatbots to exercise situational judgement. These assumptions place additional focus on what companies define as "risk" in emotionally sensitive contexts and call for companies to encode these judgments into system design. Overall, students advocate for greater corporate responsibility in the design and deployment of chatbot systems, particularly through the implementation of clearer intervention thresholds, emotional safeguards, and transparent limitations around support capabilities.

5.3 The Role of Legislation

The survey results indicate strong support for legislative involvement in determining corporate responsibility when chatbot use results in either the mishandling of mental-health–related data or negative psychological impact to a user. Across both questions regarding government involvement in data mishandling and psychological impact, a clear majority of students favored federal jurisdiction, with additional support for state-level regulation and minimal endorsement of local oversight or the absence of legislation. This pattern suggests that students perceive harms associated with chatbot systems as a systemic issue rather than isolated incidents. Students' preference for federal oversight may also reflect an understanding of the scale and reach of these technologies. Chatbots operate across state boundaries, collect and process data at scale, and are often used by corporations with national or global reach. As a result, respondents appear to view federal regulation as better positioned to establish uniform, grounding standards for accountability, data protection, and psychological harm prevention. This perception aligns similarly with recent federal legislative efforts documented in the Children Harmed by AI Technology Act, which frames psychological harm and emotional risk as consumer protection concerns enforceable through the Federal Trade Commission and state attorney general [Shah et al. 2025]. Thus, the emphasis on federal enforcement suggests that students share concerns that existing regulatory mechanisms are insufficient for governing emotionally responsive chatbot systems.

At the same time, students' support for state-level jurisdiction reflects awareness of emerging state legislation addressing chatbot use in mental-health contexts. Laws proposed or enacted in states such as New York, California, and Illinois showcase how states have begun experimenting with more targeted approaches, including mandatory self-harm detection, transparency requirements, and restriction on AI involvement in therapeutic care. Student responses suggest that while federal oversight is perceived as necessary for establishing baseline protections, state-level intervention may also be needed to establish more context-specific safeguards that respond to local priorities and variations in how chatbots are used across diverse populations. In this sense, students appear to favor a regulatory ecosystem where federal authorities provide a baseline of protection while states retain the flexibility to address gaps in federal policy and intervene more directly in mental health crises. Taken together, these findings suggest that users increasingly view emotional harm and data misuse as harms that require legal recognition and implementation through federal and state regulation.

6 Future Work

As the AI regulation space develops, corporations continue to contribute to the narratives surrounding the unmitigated concerns their products give rise to, affecting the current landscape of AI accountability. In recent months, generative AI companies have released self-reported statistics of the data they have collected of what users are using their services for, partially in an attempt to address growing concerns about unintended use. A related concern emerging from these reports is the misalignment in metrics used to indicate harm, such as message counts that pertain to therapeutic conversations, as opposed to the number of use cases that contain therapeutic counsel. Quantifying harm will be a necessary component of developing legislation in response to AI in mental health environments on any level, and an analysis of public perception of how this quantification should occur could bolster attempts to understand legislative possibility in this space. In our effort to maintain ethical standards, the survey focuses more closely on individuals' perspectives of broad statements applicable to society at large as opposed to specific and/or personal scenarios. Therefore, an unimaginable amount of nuance in that many of our questions remains overlooked, and future work can expand on the gray areas left unexplored.

7 Appendix

Survey Questions

1. Please select your age range: <18, 18-20, 21-23, 24+
2. On a scale of 1-5, how well do you understand how generative AI systems (like ChatGPT) are trained and operated? (1 = I do not understand how generative AI systems are trained/operated at all; 5 = I understand how generative AI systems are trained/operated very well)
3. For what purposes have you used generative AI chatbots? Select all that apply: For help on schoolwork, For help on work outside of school, For therapeutic support, I have never used chatbots for any purpose
4. If you have used a chatbot for advice on personal matters (ie. mental and emotional well-being), what is the strongest motivation for doing so? If "other," please specify. Select one of the following: Anonymity, Convenience, Curiousity, Emotional reflief, Lack of other support, N/A
5. If you have used a chatbot for therapeutic support, what information do you typically provide the chatbot with during your conversation? Select all that apply: Names, Ages, Locations, Educational information, Financial information, Medical information, I do not provide the chatbot with any information beyond situational context, I have not used a chatbot for therapeutic support
6. On a scale of 1-5, how much do you agree with the following statement: I know people who use chatbots for therapeutic support? (1 = I do not know anyone who uses chatbots for therapeutic support, 5 = All of the people I know who use chatbots have used it for therapeutic support)
7. On a scale of 1-5, how much do you agree with the following statement: People can trust chatbots to provide empathetic, accurate, and supportive mental health advice? (1 = I do not trust chatbots to provide empathetic, accurate, and supportive mental health advice, 5 = I trust chatbots to provide empathetic, accurate, and supportive mental health advice)
8. On a scale of 1-5, how much do you agree with the following statement: In general, people can trust chatbots with their mental health-related data and the maintenance of anonymity? (1 = Chatbots cannot be trusted with people's mental health data, 5 = Chatbots can be trusted with people's mental health data)
9. Who should be held responsible if a user is negatively emotionally impacted after interacting with a chatbot? Select one of the following. If "other," please specify: The company incurs all responsibility, The company incurs most responsibility, The user incurs most responsibility, The user incurs all responsibility, Both the company and the user bear equal responsibility

10. Who should be held responsible if sensitive user data shared with a chatbot is leaked, stolen, or misused? Select one of the following. If "other," please specify: The company incurs all responsibility, The company incurs most responsibility, The user incurs most responsibility, The user incurs all responsibility, Both the company and the user bear equal responsibility
11. Under what circumstances, if any, do you think chatbot companies should intervene during emotionally sensitive queries? Select one of the following. If "other," please specify: Only if there is a risk of harm to the user or others, Only if the user asks for help, Always intervene when emotionally sensitive engagements are detected, Never intervene
12. Should legislation (and if so, at what level) have jurisdiction over deciding a corporation's legal responsibility when mental health related data is mismanaged as a result of using their product? Select one of the following: Yes, federal government, Yes, state government, Yes, local government, No
13. Should legislation (and if so, at what level) have jurisdiction over deciding a corporation's legal responsibility when a user experiences negative psychological impact as a result of using their product? Select one of the following: Yes, federal government, Yes, state government, Yes, local government, No
14. Are you open to a 10-15 minute interview with the research team to discuss further? Your anonymity will be preserved in any report of our findings. Select one of the following: Yes, No
15. If yes, please provide your email for the research team to contact you.

References

John G. Browning. 2025. Whose bot is it anyway? Determining Liability for AI-Generated Content. *Northern Illinois University Law Review* 45, 3 (2025). https://huskiecommons.lib.niu.edu/niulr/vol45/iss3/1

Krisha Cerilli, Jessica Colnago, Julia Horwitz, Amritha Jayanti, and Stephanie T. Nguyen. 2024. *AI Companies: Uphold Your Privacy and Confidentiality Commitments.*

Mohit Chandra, Suchismita Naik, Denae Ford, Ebele Okoli, Munmun De Choudhury, Mahsa Ershadi, Gonzalo Ramos, Javier Hernandez, Ananya Bhattacharjee, Shahed Warreth, and Jina Suh. 2025. From Lived Experience to Insight: Unpacking the Psychological Risks of Using AI Conversational Agents. *FAccT '25: Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency* (2025), 975–1004. https://doi.org/10.1145/3715275.3732063

JP De Mello Barreto. 2025. *Balancing safety and privacy: regulatory models for AI misuse.* https://law-ai.org/balancing-safety-and-privacy-regulatory-models-for-ai-misuse/

Henry L. Fraser and Nicolas P. Suzor. 2025. Locating fault for AI harms: a systems theory of foreseeability, reasonable care and causal responsibility in the AI value chain. *Law Innovation and Technology* 17, 1 (2025), 103–138. https://doi.org/10.1080/17579961.2025.2469345

General Data Protection Regulation. 2016. *Article 32 - Security of Processing.* https://gdpr-info.eu/art-32-gdpr/

Kashmir Hill and Dylan Freedman. 2025. *Chatbots Can Go Into a Delusional Spiral. Here's How It Happens.* https://www.nytimes.com/2025/08/08/technology/ai-chatbots-delusions-chatgpt.html

Stylianos Kampakis. 2025. *Liability in AI-Driven Decisions: Who is Responsible When AI Gets It Wrong?* https://community.ibm.com/community/user/blogs/stylianos-kampakis/2025/06/10/liability-in-ai-driven-decisions

Jennifer King and Caroline Meinhardt. 2024. Rethinking Privacy in the AI Era: Policy Provocations for a Data-Centric World. *Stanford University Human-Centered Artificial Intelligence* (2024). https://hai.stanford.edu/policy/white-paper-rethinking-privacy-ai-era-policy-provocations-data-centric-world

J. K. C. Kingston. 2016. Artificial Intelligence and Legal Liability. *International Conference on Innovative Techniques and Applications of Artificial Intelligence* (2016), 269–279. https://doi.org/10.1007/978-3-319-47175-4_20

Jabari Kwesi, Jiaxun Cao, Riya Manchanda, and Pardis Emami-Naeini. 2025. Exploring User Security and Privacy Attitudes and Concerns Toward the Use of General-Purpose LLM Chatbots for Mental Health. *34th USENIX Security Symposium* (2025). https://www.usenix.org/conference/usenixsecurity25/presentation/kwesi

Alaap B. Shah, Frances M. Green, James A. Boiani, Eleanor T. Chung, and Epstein Becker. 2025. *Novel Lawsuits Allege AI Chatbots Encouraged Minors' Suicides, Mental Health Trauma: Considerations for Stakeholders.* https://natlawreview.com/article/novel-lawsuits-allege-ai-chatbots-encouraged-minors-suicides-mental-health-trauma

Steven Siddals, John Torous, and Astrid Coxon. 2024. 'It happened to be the perfect thing': experiences of generative AI chatbots for mental health. *Npj Mental Health Research* 3, 1 (2024), 48. https://doi.org/10.1038/s44184-024-00097-4

U.S. Food and Drug Administration. 2025. Digital Health Advisory Committee Meeting Announcement. https://www.fda.gov/advisory-committees/advisory-committee-calendar/november-6-2025-digital-health-advisory-committee-meeting-announcement-11062025#event-information

Aditya Nrusimha Vaidyam, Hannah Wisniewski, John David Halamka, Matcheri S Kashavan, and John Blake Torous. 2019. Chatbots and Conversational Agents in Mental Health: A Review of the Psychiatric Landscape. *Can J Psychiatry* 64, 7 (2019). https://doi.org/10.1177/0706743719828977

Woebot Health. 2021. *New Survey Shows Paradigm Shift in American Attitudes and Opinions About Mental Health Technology.* https://woebothealth.com/new-survey-shows-paradigm-shift-in-american-attitudes-and-opinions-about-mental-health-technology/

Marc Zao-Sanders. 2025. *How People Are Really Using Gen AI in 2025.* https://hbr.org/2025/04/how-people-are-really-using-gen-ai-in-2025

Paulius Čerka, Jurgita Grigienė, and Gintarė Sirbikytė. 2015. Liability for damages caused by artificial intelligence. *Computer Law and Security Review* 31, 3 (2015), 376–389. https://doi.org/10.1016/j.clsr.2015.03.008